

You have 2 free stories left this month. Sign up and get an extra one for free.

Building a Data API with FastAPI and SQLAlchemy

How to use modular database logic to load data into a database and perform CRUD operations in an App



Edward Krueger

May 8 · 6 min read ★

By: Edward Krueger Data Scientist and Instructor and Douglas Franklin Teaching Assistant and Technical Writer.

Link to Github repo with app code: <https://github.com/edkrueger/sars-fastapi>





Photo by Chris Liverani on Unsplash

What is FastAPI?

FastAPI is a high-performance API based on Pydantic and Starlette. FastAPI integrates well with many packages, including many ORMs. With FastAPI, you can use most relational databases. FastAPI easily integrates with SQLAlchemy and SQLAlchemy supports PostgreSQL, MySQL, SQLite, Oracle, Microsoft SQL Server and others.

Other python microservice frameworks like Flask don't integrate with SQLAlchemy easily. It is common to use Flask with a package called Flask-SQLAlchemy. Flask-SQLAlchemy isn't necessary and has problems of its own. For more information on this, give this article a read!

There is no FastAPI-SQLAlchemy because FastAPI integrates well with vanilla SQLAlchemy!

Modular App Structure

When writing an app, it is best to create independent and modular python code. Here we will discuss the following constituent files of our app, `database.py`, `models.py`, `schemas.py`, `main.py`, and `load.py`.

Ideally, you should only have to define your database models once! Using SQLAlchemy's `declarative_base()` and `Base.metadata.create_all()` allows you to write just one class per table to use in the app, to use in Python outside of the app and to use in the database. With a separate **`database.py`** and **`models.py`** file, we establish our database table classes and connection a single time, then call them later as needed.

To avoid confusion between the SQLAlchemy *models* and the Pydantic *models*, we will have the file `models.py` with the SQLAlchemy models, and the file `schemas.py` with the Pydantic models. Also of note is that SQLAlchemy and Pydantic use slightly different syntax to define models, as seen in the below files.

Database.py

Here is the file that defines our database connection using SQLAlchemy.

```
1  import os
2
```

```

3  from sqlalchemy import create_engine
4  from sqlalchemy.ext.declarative import declarative_base
5  from sqlalchemy.orm import sessionmaker
6
7  SQLALCHEMY_DATABASE_URL = os.getenv("DB_CONN")
8
9  engine = create_engine(SQLALCHEMY_DATABASE_URL)
10 SessionLocal = sessionmaker(autocommit=False, autoflush=False, bind=engine)
11
12 Base = declarative_base()

```

database.py hosted with ❤ by GitHub

[view raw](#)

database.py

Declarative Base and MetaData

The `declarative_base()` base class contains a `MetaData` object where newly defined `Table` objects are collected. This `MetaData` object is accessed when we call the line `models.Base.metadata.create_all()` to create all of our tables.

Session Local: Handling Threading Issues

SQLAlchemy includes a helper object that helps with the establishment of user-defined `Session` scopes. This is useful for eliminating threading issues across your app.

To create a session, below we use the `sessionmaker` function and pass it a few arguments. `Sessionmaker` is a factory for initializing new `Session` objects. `Sessionmaker` initializes these sessions by requesting a connection from the engine's connection pool and attaching a connection to the new `Session` object.

```
SessionLocal = sessionmaker(autocommit=False, autoflush=False, bind=engine)
```

SessionLocal from database.py

Initializing a new session object is also referred to as “checking out” a connection. The database stores a list of these connections/processes. So when you begin a new session, be mindful you are also starting a new process within the database. If the database doesn't have these connections closed, there is a maximum number of connections that can be reached. The database will eventually kill idle processes like stale connections; however, it can take hours before that happens.

SQLAlchemy has some pool options to prevent this, but removing the connections when they are no longer needed is best! The FastAPI docs include a `get_db()` function

that allows a route to use the same session through a request and then close it when the request is finished. Then `get_db()` creates a new session for the next request.

Once we have our database connection and session set up, we are ready to build our other app components.

Models.py

Notice that we import the `Base` class, defined in the `database.py` file above, into the `models.py` file below to use `declarative_base()`.

```
1  from sqlalchemy import Column, Integer, String
2  from sqlalchemy.types import Date
3  from .database import Base
4
5
6  class Record(Base):
7      __tablename__ = "Records"
8
9      id = Column(Integer, primary_key=True, index=True)
10     date = Column(Date)
11     country = Column(String(255), index=True)
12     cases = Column(Integer)
13     deaths = Column(Integer)
14     recoveries = Column(Integer)
```

models.py hosted with ❤ by GitHub

[view raw](#)

models.py

This file creates the model or schema for the table `Records` in our database.

Using SQLAlchemy's `declarative_base()` allows you to write just one model for each table that app uses. That model is then used in Python outside of the app and in the database.

Having these separate Python files is good because you can use the same model to query or load data outside of an app. Additionally, you'll have one version of each model, which simplifies development.

These modular Python files can be used to reference the same models or databases in data pipelines, report generation, and anywhere else they are needed.

Schemas.py

Here we write our schema for Pydantic. Remember that FastAPI is built upon Pydantic. The primary means of defining objects in Pydantic is via models that inherit from `BaseModel`.

Pydantic guarantees that the data fields of the resultant model conform to the field types we have defined, using standard modern Python types, for the model.

```
1  from datetime import date
2  from pydantic import BaseModel
3
4
5  class Record(BaseModel):
6      id: int
7      date: date
8      country: str
9      cases: int
10     deaths: int
11     recoveries: int
12
13     class Config:
14         orm_mode = True
```

schemas.py hosted with ❤ by GitHub

[view raw](#)

schemas.py

The line `orm_mode = True` allows the app to take ORM objects and translate them into responses automatically. This automation saves us from manually taking data out of ORM, making it into a dictionary, then loading it in with Pydantic.

Main.py

Here is where we bring all the modular components together.

After importing all of our dependencies and modular app components we call `models.Base.metadata.create_all(bind=engine)` to create our models in the database.

Next, we define our app and set up CORS Middleware.

```
1  from typing import List
2
3  from fastapi import Depends, FastAPI, HTTPException
4  from fastapi.middleware.cors import CORSMiddleware
5  from sqlalchemy.orm import Session
6  from starlette.responses import RedirectResponse
```

```

6 from starlette.responses import RedirectResponse
7
8 from . import models, schemas
9 from .database import SessionLocal, engine
10
11 models.Base.metadata.create_all(bind=engine)
12
13 app = FastAPI()
14
15
16 app.add_middleware(
17     CORSMiddleware,
18     allow_origins=["*"],
19     allow_methods=["*"],
20     allow_headers=["*"],
21     allow_credentials=True,
22 )
23
24 # Dependency
25 def get_db():
26     try:
27         db = SessionLocal()
28         yield db
29     finally:
30         db.close()
31
32
33 @app.get("/")
34 def main():
35     return RedirectResponse(url="/docs/")
36
37
38 @app.get("/records/", response_model=List[schemas.Record])
39 def show_records(db: Session = Depends(get_db)):
40     records = db.query(models.Record).all()
41     return records

```

CORS Middleware

CORS or “Cross-Origin Resource Sharing” refers to the situations where a frontend running in a browser has JavaScript code that communicates with a backend, and the backend is in a different “origin” than the frontend.

To configure CORS middleware in your **FastAPI** application

- Import `CORSMiddleware` .
- Create a list of allowed origins as strings, i.e., “http://localhost,” “http://localhost:8080.”
- Add it as a “middleware” to your **FastAPI** application.

You can also specify if your backend allows:

- Credentials (Authorization headers, Cookies, etc.).
- Specific HTTP methods (`POST` , `PUT`) or all of them with the wildcard `"*"` .
- Specific HTTP headers or all of them with the wildcard `"*"` .

For everything to work correctly, it’s best to specify the allowed origins explicitly. Additionally, for security reasons, it is good practice to be explicit about which origins may access your app.

Properly setting up the CORS middleware will eliminate any CORS issues within your app.

The `get_db()` function ensures that any route passed this function ought to have our `SessionLocal` database connection when needed and that the session is closed after use.

The `/records` route is for viewing our app’s data. Notice that we use `schemas.py` for the frontend and `models.py` to query our backend in this route.

```
1 @app.get("/records/", response_model=List[schemas.Record])
2 def show_records(db: Session = Depends(get_db)):
3     records = db.query(models.Record).all()
4     return records
```

main.py hosted with ❤ by GitHub

[view raw](#)

FastAPI /records route

External Data Loading

Instead of using the app to load data, we load the database with a separate Python file.

Here is an example Python file that reads data from a CSV and inserts that data into a database.

```
1  import csv
2  import datetime
3
4  from app import models
5  from app.database import SessionLocal, engine
6
7  db = SessionLocal()
8
9  models.Base.metadata.create_all(bind=engine)
10
11 with open("sars_2003_complete_dataset_clean.csv", "r") as f:
12     csv_reader = csv.DictReader(f)
13
14     for row in csv_reader:
15         db_record = models.Record(
16             date=datetime.datetime.strptime(row["date"], "%Y-%m-%d"),
17             country=row["country"],
18             cases=row["cases"],
19             deaths=row["deaths"],
20             recoveries=row["recoveries"],
21         )
22         db.add(db_record)
23
24     db.commit()
25
26 db.close()
```

load.py hosted with  by GitHub

[view raw](#)

load.py

Notice that we import the models, our custom session `SessionLocal`, and our engine that we've defined in other Python files. Then we read the CSV and using the `models.Record` schema, add `db_record` to the database through the `SessionLocal()` connection.

Database Loading

If your app is set up properly, including your database connection string you may call:

```
python load.py
```


This will load your local or remote database without ever having to run the app!

Testing the app

To run the app locally using a remote DB, in the terminal run:

```
uvicorn app.main:app
```

This runs your app locally. This local instance is connected to the cloud database we just loaded, so check the /records route to see the data!

Conclusion

The modular app structure used allows us to define a model once then use it as needed. This practice makes development much easier because you only have one file to debug if something goes awry. Additionally, your code will be much more reusable and ready for another project!

For your next data project, whether it is a dashboard, data journal, or an API, be sure to give FastAPI a try! It is very fast and easy to pick up, and they have excellent documentation to guide you along the way.

We hope this has been helpful, thank you for reading, and best of luck!

Sign up for The Daily Pick

By Towards Data Science

Hands-on real-world examples, research, tutorials, and cutting-edge techniques delivered Monday to Thursday. Make learning your daily ritual. [Take a look](#)

Your email

Get this newsletter

By signing up, you will create a Medium account if you don't already have one. Review our [Privacy Policy](#) for more information about our privacy practices.

[API](#)[Database](#)[Software Development](#)[Sqlalchemy](#)[Apps](#)[About](#) [Help](#) [Legal](#)

Get the Medium app

